

```
!pip install pafy youtube-dl moviepy
!pip install imageio-ffmpeg
!pip3 install imageio==2.4.1
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: pafy in /usr/local/lib/python3.7/dist-packages (0.5.5)
Requirement already satisfied: youtube-dl in /usr/local/lib/python3.7/dist-packages (2021.12.1)
Requirement already satisfied: moviepy in /usr/local/lib/python3.7/dist-packages (0.2.3)
Requirement already satisfied: decorator<5.0,>=4.0.2 in /usr/local/lib/python3.7/dist-packages (4.4.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from moviepy) (1.21.0)
Requirement already satisfied: tqdm<5.0,>=4.11.2 in /usr/local/lib/python3.7/dist-packages (from moviepy) (4.62.3)
Requirement already satisfied: imageio<3.0,>=2.1.2 in /usr/local/lib/python3.7/dist-packages (from imageio-ffmpeg) (2.14.1)
Requirement already satisfied: pillow in /usr/local/lib/python3.7/dist-packages (from imageio) (8.3.2)
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: imageio-ffmpeg in /usr/local/lib/python3.7/dist-packages (0.0.4)
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: imageio==2.4.1 in /usr/local/lib/python3.7/dist-packages (2.4.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from imageio) (1.21.0)
Requirement already satisfied: pillow in /usr/local/lib/python3.7/dist-packages (from imageio) (8.3.2)
```

```
# Import the required libraries.
```

```
import os
import cv2
import pafy
import math
import random
import numpy as np
import datetime as dt
import tensorflow as tf
from collections import deque
import matplotlib.pyplot as plt
```

```
from moviepy.editor import *
%matplotlib inline
```

```
from sklearn.model_selection import train_test_split
```

```
from tensorflow.keras.layers import *
from tensorflow.keras.models import Sequential
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.utils import plot_model
```

```
seed_constant = 27
np.random.seed(seed_constant)
random.seed(seed_constant)
tf.random.set_seed(seed_constant)
```

```

# Discard the output of this cell.
%%capture

# Downlaod the UCF50 Dataset
!wget --no-check-certificate https://www.crcv.ucf.edu/data/UCF50.rar

#Extract the Dataset
!unrar x UCF50.rar


plt.figure(figsize = (20, 20))

# Get the names of all classes/categories in UCF50.
all_classes_names = os.listdir('UCF50')

# Generate a list of 20 random values. The values will be between 0-50,
# where 50 is the total number of class in the dataset.
random_range = random.sample(range(len(all_classes_names)), 20)

# Iterating through all the generated random values.
for counter, random_index in enumerate(random_range, 1):

    # Retrieve a Class Name using the Random Index.
    selected_class_Name = all_classes_names[random_index]

    # Retrieve the list of all the video files present in the randomly selected Class Directo
    video_files_names_list = os.listdir(f'UCF50/{selected_class_Name}')

    # Randomly select a video file from the list retrieved from the randomly selected Class D
    selected_video_file_name = random.choice(video_files_names_list)

    # Initialize a VideoCapture object to read from the video File.
    video_reader = cv2.VideoCapture(f'UCF50/{selected_class_Name}/{selected_video_file_name}')

    # Read the first frame of the video file.
    _, bgr_frame = video_reader.read()






    # Release the VideoCapture object.
    video_reader.release()

    # Convert the frame from BGR into RGB format.
    rgb_frame = cv2.cvtColor(bgr_frame, cv2.COLOR_BGR2RGB)

    # Write the class name on the video frame.
    cv2.putText(rgb_frame, selected_class_Name, (10, 30), cv2.FONT_HERSHEY_SIMPLEX, 1, (255,

# Display the frame.
plt.subplot(5, 4, counter);plt.imshow(rgb_frame);plt.axis('off')

```



```

# Specify the height and width to which each video frame will be resized in our dataset.
IMAGE_HEIGHT , IMAGE_WIDTH = 64, 64

# Specify the number of frames of a video that will be fed to the model as one sequence.
SEQUENCE_LENGTH = 20

# Specify the directory containing the UCF50 dataset.
DATASET_DIR = "UCF50"

# Specify the list containing the names of the classes used for training. Feel free to choose
CLASSES_LIST = ["WalkingWithDog", "TaiChi", "Swing", "HorseRace"]

def frames_extraction(video_path):
    """
    This function will extract the required frames from a video after resizing and normalizing.
    Args:
        video_path: The path of the video in the disk, whose frames are to be extracted.
    Returns:
        frames_list: A list containing the resized and normalized frames of the video.
    """

    # Declare a list to store video frames.
    frames_list = []

    # Read the Video File using the VideoCapture object.
    video_reader = cv2.VideoCapture(video_path)

    # Get the total number of frames in the video.
    video_frames_count = int(video_reader.get(cv2.CAP_PROP_FRAME_COUNT))

    # Calculate the interval after which frames will be added to the list.
    skip_frames_window = max(int(video_frames_count/SEQUENCE_LENGTH), 1)

    # Iterate through the Video Frames.
    for frame_counter in range(SEQUENCE_LENGTH):

        # Set the current frame position of the video.
        video_reader.set(cv2.CAP_PROP_POS_FRAMES, frame_counter * skip_frames_window)

        # Reading the frame from the video.
        success, frame = video_reader.read()

        # Check if Video frame is not successfully read then break the loop
        if not success:
            break

        # Resize the Frame to fixed height and width.
        resized_frame = cv2.resize(frame, (IMAGE_HEIGHT, IMAGE_WIDTH))

        # Normalize the resized frame by dividing it with 255 so that each pixel value then 1

```

```

        normalized_frame = resized_frame / 255

        # Append the normalized frame into the frames list
        frames_list.append(normalized_frame)

    # Release the VideoCapture object.
    video_reader.release()

    # Return the frames list.
    return frames_list

def create_dataset():
    """
    This function will extract the data of the selected classes and create the required datas
    Returns:
        features:          A list containing the extracted frames of the videos.
        labels:            A list containing the indexes of the classes associated with the v
        video_files_paths: A list containing the paths of the videos in the disk.
    """

    # Declared Empty Lists to store the features, labels and video file path values.
    features = []
    labels = []
    video_files_paths = []

    # Iterating through all the classes mentioned in the classes list
    for class_index, class_name in enumerate(CLASSES_LIST):

        # Display the name of the class whose data is being extracted.
        print(f'Extracting Data of Class: {class_name}')

        # Get the list of video files present in the specific class name directory.
        files_list = os.listdir(os.path.join(DATASET_DIR, class_name))

        # Iterate through all the files present in the files list.
        for file_name in files_list:

            # Get the complete video path.
            video_file_path = os.path.join(DATASET_DIR, class_name, file_name)

            # Extract the frames of the video file.
            frames = frames_extraction(video_file_path)

            # Check if the extracted frames are equal to the SEQUENCE_LENGTH specified above.
            # So ignore the vides having frames less than the SEQUENCE_LENGTH.
            if len(frames) == SEQUENCE_LENGTH:

                # Append the data to their repective lists.
                features.append(frames)
                labels.append(class_index)

```

```

        video_files_paths.append(video_file_path)

# Converting the list to numpy arrays
features = np.asarray(features)
labels = np.array(labels)

# Return the frames, class index, and video file path.
return features, labels, video_files_paths

# Create the dataset.
features, labels, video_files_paths = create_dataset()

    Extracting Data of Class: WalkingWithDog
    Extracting Data of Class: TaiChi
    Extracting Data of Class: Swing
    Extracting Data of Class: HorseRace

1
2
# Using Keras's to_categorical method to convert labels into one-hot-encoded vectors
one_hot_encoded_labels = to_categorical(labels)

def create_convlstm_model():
    '''
    This function will construct the required convlstm model.
    Returns:
        model: It is the required constructed convlstm model.
    '''

    # We will use a Sequential model for model construction
    model = Sequential()

    # Define the Model Architecture.
    #####

    model.add(ConvLSTM2D(filters = 4, kernel_size = (3, 3), activation = 'tanh', data_format =
        recurrent_dropout=0.2, return_sequences=True, input_shape = (SEQUENC
        IMAGE_H

    model.add(MaxPooling3D(pool_size=(1, 2, 2), padding='same', data_format='channels_last'))
    model.add(TimeDistributed(Dropout(0.2)))

    model.add(ConvLSTM2D(filters = 8, kernel_size = (3, 3), activation = 'tanh', data_format
        recurrent_dropout=0.2, return_sequences=True))

    model.add(MaxPooling3D(pool_size=(1, 2, 2), padding='same', data_format='channels_last'))
    model.add(TimeDistributed(Dropout(0.2)))

```

```
model.add(ConvLSTM2D(filters = 14, kernel_size = (3, 3), activation = 'tanh', data_format
                    recurrent_dropout=0.2, return_sequences=True))

model.add(MaxPooling3D(pool_size=(1, 2, 2), padding='same', data_format='channels_last'))
model.add(TimeDistributed(Dropout(0.2)))

model.add(ConvLSTM2D(filters = 16, kernel_size = (3, 3), activation = 'tanh', data_format
                    recurrent_dropout=0.2, return_sequences=True))

model.add(MaxPooling3D(pool_size=(1, 2, 2), padding='same', data_format='channels_last'))
#model.add(TimeDistributed(Dropout(0.2)))

model.add(Flatten())

model.add(Dense(len(CLASSES_LIST), activation = "softmax"))

#####

# Display the models summary.
model.summary()

# Return the constructed convlstm model.
return model
```

```
# Construct the required convlstm model.
convlstm_model = create_convlstm_model()

# Display the success message.
print("Model Created Successfully!")
```

Model: "sequential"

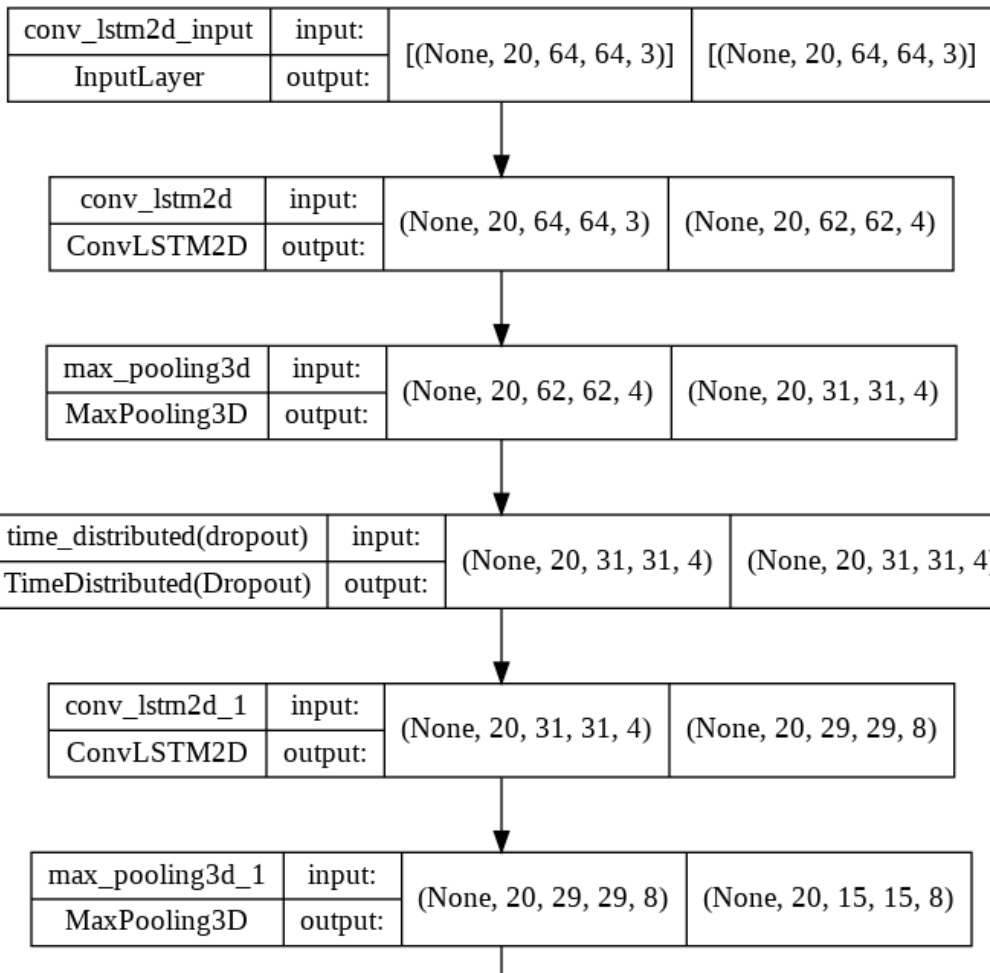
Layer (type)	Output Shape	Param #
=====		
conv_lstm2d (ConvLSTM2D)	(None, 20, 62, 62, 4)	1024
max_pooling3d (MaxPooling3D)	(None, 20, 31, 31, 4)	0
time_distributed (TimeDistributed)	(None, 20, 31, 31, 4)	0
conv_lstm2d_1 (ConvLSTM2D)	(None, 20, 29, 29, 8)	3488
max_pooling3d_1 (MaxPooling3D)	(None, 20, 15, 15, 8)	0
time_distributed_1 (TimeDistributed)	(None, 20, 15, 15, 8)	0
conv_lstm2d_2 (ConvLSTM2D)	(None, 20, 13, 13, 14)	11144

max_pooling3d_2 (MaxPooling 3D)	(None, 20, 7, 7, 14)	0
time_distributed_2 (TimeDistributed)	(None, 20, 7, 7, 14)	0
conv_lstm2d_3 (ConvLSTM2D)	(None, 20, 5, 5, 16)	17344
max_pooling3d_3 (MaxPooling 3D)	(None, 20, 3, 3, 16)	0
flatten (Flatten)	(None, 2880)	0
dense (Dense)	(None, 4)	11524

```
=====
Total params: 44,524
Trainable params: 44,524
Non-trainable params: 0
```

Model Created Successfully!

```
# Plot the structure of the constructed model.
plot_model(convlstm_model, to_file = 'convlstm_model_structure_plot.png', show_shapes = True,
```

```

1
2
# Split the Data into Train ( 75% ) and Test Set ( 25% ).
features_train, features_test, labels_train, labels_test = train_test_split(features, one_hot.
    | conv_lstm2d_2 | input: | (None, 20, 45, 45, 8) | (None, 20, 43, 43, 8) |

# Create an Instance of Early Stopping Callback
early_stopping_callback = EarlyStopping(monitor = 'val_loss', patience = 10, mode = 'min', re

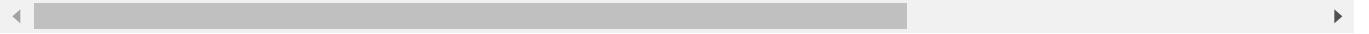
# Compile the model and specify loss function, optimizer and metrics values to the model
convlstm_model.compile(loss = 'categorical_crossentropy', optimizer = 'Adam', metrics = ["acc

# Start training the model.
convlstm_model_training_history = convlstm_model.fit(x = features_train, y = labels_train, ep

Epoch 1/50
73/73 [=====] - 148s 2s/step - loss: 1.3874 - accuracy: 0.2979
Epoch 2/50
73/73 [=====] - 138s 2s/step - loss: 1.3407 - accuracy: 0.3322
Epoch 3/50
73/73 [=====] - 139s 2s/step - loss: 1.2268 - accuracy: 0.4760
Epoch 4/50
73/73 [=====] - 139s 2s/step - loss: 0.9917 - accuracy: 0.6096
Epoch 5/50

```

```
73/73 [=====] - 139s 2s/step - loss: 0.7571 - accuracy: 0.6918
Epoch 6/50
73/73 [=====] - 139s 2s/step - loss: 0.6136 - accuracy: 0.7534
Epoch 7/50
73/73 [=====] - 139s 2s/step - loss: 0.5121 - accuracy: 0.7979
Epoch 8/50
73/73 [=====] - 139s 2s/step - loss: 0.3390 - accuracy: 0.8699
Epoch 9/50
73/73 [=====] - 140s 2s/step - loss: 0.2729 - accuracy: 0.9007
Epoch 10/50
73/73 [=====] - 138s 2s/step - loss: 0.2915 - accuracy: 0.8801
Epoch 11/50
73/73 [=====] - 138s 2s/step - loss: 0.2309 - accuracy: 0.9041
Epoch 12/50
73/73 [=====] - 139s 2s/step - loss: 0.0646 - accuracy: 0.9863
Epoch 13/50
73/73 [=====] - 139s 2s/step - loss: 0.0747 - accuracy: 0.9760
Epoch 14/50
73/73 [=====] - 139s 2s/step - loss: 0.0812 - accuracy: 0.9692
Epoch 15/50
73/73 [=====] - 139s 2s/step - loss: 0.0682 - accuracy: 0.9863
Epoch 16/50
73/73 [=====] - 139s 2s/step - loss: 0.0914 - accuracy: 0.9658
Epoch 17/50
73/73 [=====] - 139s 2s/step - loss: 0.0479 - accuracy: 0.9897
Epoch 18/50
73/73 [=====] - 138s 2s/step - loss: 0.0147 - accuracy: 1.0000
Epoch 19/50
73/73 [=====] - 139s 2s/step - loss: 0.0228 - accuracy: 0.9966
Epoch 20/50
73/73 [=====] - 138s 2s/step - loss: 0.1182 - accuracy: 0.9623
Epoch 21/50
73/73 [=====] - 138s 2s/step - loss: 0.0846 - accuracy: 0.9795
Epoch 22/50
73/73 [=====] - 138s 2s/step - loss: 0.0529 - accuracy: 0.9760
Epoch 23/50
73/73 [=====] - 138s 2s/step - loss: 0.0284 - accuracy: 0.9897
```



```
# Evaluate the trained model.
```

```
model_evaluation_history = convlstm_model.evaluate(features_test, labels_test)
```

```
4/4 [=====] - 14s 3s/step - loss: 0.7390 - accuracy: 0.8033
```

```
# Get the loss and accuracy from model_evaluation_history.
```

```
model_evaluation_loss, model_evaluation_accuracy = model_evaluation_history
```

```
# Define the string date format.
```

```
# Get the current Date and Time in a DateTime Object.
```

```
# Convert the DateTime object to string according to the style mentioned in date_time_format
```

```

date_time_format = '%Y_%m_%d__%H_%M_%S'
current_date_time_dt = dt.datetime.now()
current_date_time_string = dt.datetime.strftime(current_date_time_dt, date_time_format)

# Define a useful name for our model to make it easy for us while navigating through multiple
model_file_name = f'convlstm_model__Date_Time_{current_date_time_string}__Loss_{model_evalu

# Save your Model.
convlstm_model.save(model_file_name)

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
def plot_metric(model_training_history, metric_name_1, metric_name_2, plot_name):
    ...

    This function will plot the metrics passed to it in a graph.
    Args:
        model_training_history: A history object containing a record of training and validati
                                loss values and metrics values at successive epochs
        metric_name_1:          The name of the first metric that needs to be plotted in the
        metric_name_2:          The name of the second metric that needs to be plotted in the
        plot_name:              The title of the graph.
    ...

    # Get metric values using metric names as identifiers.
    metric_value_1 = model_training_history.history[metric_name_1]

```

```

metric_value_2 = model_training_history.history[metric_name_2]

# Construct a range object which will be used as x-axis (horizontal plane) of the graph.
epochs = range(len(metric_value_1))

# Plot the Graph.
plt.plot(epochs, metric_value_1, 'blue', label = metric_name_1)
plt.plot(epochs, metric_value_2, 'red', label = metric_name_2)

# Add title to the plot.
plt.title(str(plot_name))


# Add legend to the plot.
plt.legend()

```

```

# Visualize the training and validation loss metrics.
plot_metric(convlstm_model_training_history, 'loss', 'val_loss', 'Total Loss vs Total Validat

```



```

# Visualize the training and validation accuracy metrics.
plot_metric(convlstm_model_training_history, 'accuracy', 'val_accuracy', 'Total Accuracy vs T

```



- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11
- 12
- 13
- 14
- 15
- 16
- 17
- 18
- 19
- 20
- 21
- 22
- 23
- 24
- 25
- 26
- 27
- 28
- 29
- 30
- 31
- 32
- 33
- 34
- 35
- 36
- 37
- 38
- 39
- 40
- 41
- 42
- 43
- 44

```

def create_LRCN_model():
    ...

    This function will construct the required LRCN model.
    Returns:
        model: It is the required constructed LRCN model.
    ...

    # We will use a Sequential model for model construction.
    model = Sequential()

    # Define the Model Architecture.
    #####

    model.add(TimeDistributed(Conv2D(16, (3, 3), padding='same', activation = 'relu'),
                               input_shape = (SEQUENCE_LENGTH, IMAGE_HEIGHT, IMAGE_WIDTH, 3)))

    model.add(TimeDistributed(MaxPooling2D((4, 4))))
    model.add(TimeDistributed(Dropout(0.25)))

    model.add(TimeDistributed(Conv2D(32, (3, 3), padding='same', activation = 'relu')))
    model.add(TimeDistributed(MaxPooling2D((4, 4))))
    model.add(TimeDistributed(Dropout(0.25)))

    model.add(TimeDistributed(Conv2D(64, (3, 3), padding='same', activation = 'relu')))
    model.add(TimeDistributed(MaxPooling2D((2, 2))))
    model.add(TimeDistributed(Dropout(0.25)))

    model.add(TimeDistributed(Conv2D(64, (3, 3), padding='same', activation = 'relu')))
    model.add(TimeDistributed(MaxPooling2D((2, 2))))
    #model.add(TimeDistributed(Dropout(0.25)))

    model.add(TimeDistributed(Flatten()))

    model.add(LSTM(32))

    model.add(Dense(len(CLASSES_LIST), activation = 'softmax'))

    #####

    # Display the models summary.
    model.summary()

    # Return the constructed LRCN model.
    return model

# Construct the required LRCN model.
LRCN_model = create_LRCN_model()

# Display the success message.
print("Model Created Successfully!")

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
time_distributed_3 (TimeDistributed)	(None, 20, 64, 64, 16)	448
time_distributed_4 (TimeDistributed)	(None, 20, 16, 16, 16)	0
time_distributed_5 (TimeDistributed)	(None, 20, 16, 16, 16)	0
time_distributed_6 (TimeDistributed)	(None, 20, 16, 16, 32)	4640
time_distributed_7 (TimeDistributed)	(None, 20, 4, 4, 32)	0
time_distributed_8 (TimeDistributed)	(None, 20, 4, 4, 32)	0
time_distributed_9 (TimeDistributed)	(None, 20, 4, 4, 64)	18496
time_distributed_10 (TimeDistributed)	(None, 20, 2, 2, 64)	0
time_distributed_11 (TimeDistributed)	(None, 20, 2, 2, 64)	0
time_distributed_12 (TimeDistributed)	(None, 20, 2, 2, 64)	36928
time_distributed_13 (TimeDistributed)	(None, 20, 1, 1, 64)	0
time_distributed_14 (TimeDistributed)	(None, 20, 64)	0
lstm (LSTM)	(None, 32)	12416
dense_1 (Dense)	(None, 4)	132
=====		
Total params: 73,060		
Trainable params: 73,060		
Non-trainable params: 0		
=====		
Model Created Successfully!		

```
# Plot the structure of the constructed LRCN model.  
plot_model(LRCN_model, to_file = 'LRCN_model_structure_plot.png', show_shapes = True, show_la
```


time_distributed_3_input	input:	[(None, 20, 64, 64, 3)]	[(None, 20, 64, 64, 3)]
InputLayer	output:		



```

1
2
3
4
5
6
7
8
# Create an Instance of Early Stopping Callback.
early_stopping_callback = EarlyStopping(monitor = 'val_loss', patience = 15, mode = 'min', re

# Compile the model and specify loss function, optimizer and metrics to the model.
LRCN_model.compile(loss = 'categorical_crossentropy', optimizer = 'Adam', metrics = ["accurac

# Start training the model.
LRCN_model_training_history = LRCN_model.fit(x = features_train, y = labels_train, epochs = 5


Epoch 1/50
73/73 [=====] - 16s 188ms/step - loss: 1.4006 - accuracy: 0.274
Epoch 2/50
73/73 [=====] - 13s 177ms/step - loss: 1.3551 - accuracy: 0.291
Epoch 3/50
73/73 [=====] - 13s 178ms/step - loss: 1.2244 - accuracy: 0.455
Epoch 4/50
73/73 [=====] - 13s 178ms/step - loss: 1.1435 - accuracy: 0.516
Epoch 5/50
73/73 [=====] - 13s 179ms/step - loss: 1.0011 - accuracy: 0.602
Epoch 6/50
73/73 [=====] - 13s 178ms/step - loss: 0.9180 - accuracy: 0.616
Epoch 7/50
73/73 [=====] - 13s 178ms/step - loss: 0.9012 - accuracy: 0.616
Epoch 8/50
73/73 [=====] - 13s 177ms/step - loss: 0.9127 - accuracy: 0.661
Epoch 9/50
73/73 [=====] - 13s 178ms/step - loss: 0.7579 - accuracy: 0.688
Epoch 10/50
73/73 [=====] - 13s 178ms/step - loss: 0.6774 - accuracy: 0.732
Epoch 11/50
73/73 [=====] - 13s 177ms/step - loss: 0.5738 - accuracy: 0.787
Epoch 12/50
73/73 [=====] - 13s 178ms/step - loss: 0.5762 - accuracy: 0.791
Epoch 13/50
73/73 [=====] - 13s 176ms/step - loss: 0.5474 - accuracy: 0.791
Epoch 14/50
73/73 [=====] - 13s 177ms/step - loss: 0.3842 - accuracy: 0.875
Epoch 15/50
73/73 [=====] - 13s 177ms/step - loss: 0.3665 - accuracy: 0.869
Epoch 16/50

```

```

73/73 [=====] - 13s 177ms/step - loss: 0.3413 - accuracy: 0.887
Epoch 17/50
73/73 [=====] - 13s 178ms/step - loss: 0.2634 - accuracy: 0.907
Epoch 18/50
73/73 [=====] - 13s 177ms/step - loss: 0.2320 - accuracy: 0.928
Epoch 19/50
73/73 [=====] - 13s 178ms/step - loss: 0.2178 - accuracy: 0.938
Epoch 20/50
73/73 [=====] - 13s 178ms/step - loss: 0.3693 - accuracy: 0.886
Epoch 21/50
73/73 [=====] - 13s 178ms/step - loss: 0.1706 - accuracy: 0.948
Epoch 22/50
55/73 [=====>.....] - ETA: 2s - loss: 0.2385 - accuracy: 0.9227

```



```

1
2
# Evaluate the trained model.
model_evaluation_history = LRCN_model.evaluate(features_test, labels_test)

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
# Get the loss and accuracy from model_evaluation_history.
model_evaluation_loss, model_evaluation_accuracy = model_evaluation_history

# Define the string date format.
# Get the current Date and Time in a DateTime Object.
# Convert the DateTime object to string according to the style mentioned in date_time_format
date_time_format = '%Y_%m_%d_%H_%M_%S'
current_date_time_dt = dt.datetime.now()
current_date_time_string = dt.datetime.strftime(current_date_time_dt, date_time_format)

# Define a useful name for our model to make it easy for us while navigating through multiple
model_file_name = f'LRCN_model__Date_Time_{current_date_time_string}__Loss_{model_evaluation_accuracy}'

```

```
# Save the Model.  
LRCN_model.save(model_file_name)
```

```
1  
2  
# Visualize the training and validation loss metrics.  
plot_metric(LRCN_model_training_history, 'loss', 'val_loss', 'Total Loss vs Total Validation
```

```
1  
2  
# Visualize the training and validation accuracy metrics.  
plot_metric(LRCN_model_training_history, 'accuracy', 'val_accuracy', 'Total Accuracy vs Total
```

```
1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
def download_youtube_videos(youtube_video_url, output_directory):  
    '''  
    This function downloads the youtube video whose URL is passed to it as an argument.  
    Args:  
        youtube_video_url: URL of the video that is required to be downloaded.  
        output_directory: The directory path to which the video needs to be stored after dow
```

Returns:

```
    title: The title of the downloaded youtube video.  
    ...
```

```
# Create a video object which contains useful information about the video.  
video = pafy.new(youtube_video_url)
```

```
# Retrieve the title of the video.  
title = video.title
```

```
# Get the best available quality object for the video.  
video_best = video.getbest()
```

```
# Construct the output file path.  
output_file_path = f'{output_directory}/{title}.mp4'
```

```
# Download the youtube video at the best available quality and store it to the contructe  
video_best.download(filepath = output_file_path, quiet = True)
```

```
# Return the video title.  
return title
```

```
!pip install git+https://github.com/Cupcakus/pafy
```

```
!pip uninstall -y pafy
```

```
!pip install git+https://github.com/Cupcakus/pafy
```

```
1  
2  
3  
4  
5  
6  
7  
8  
9
```

```
# Make the Output directory if it does not exist  
test_videos_directory = 'test_videos'  
os.makedirs(test_videos_directory, exist_ok = True)
```

```
# Download a YouTube Video.  
video_title = download_youtube_videos('https://www.youtube.com/watch?v=8u0qjmHIOcE', test_vid
```

```
# Get the YouTube Video's path we just downloaded.  
input_video_file_path = f'{test_videos_directory}/{video_title}.mp4'
```

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50

```

51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
def predict_on_video(video_file_path, output_file_path, SEQUENCE_LENGTH):
    '''
    This function will perform action recognition on a video using the LRCN model.
    Args:
    video_file_path: The path of the video stored in the disk on which the action recognitio
    output_file_path: The path where the ouput video with the predicted action being performe
    SEQUENCE_LENGTH: The fixed number of frames of a video that can be passed to the model a
    '''

    # Initialize the VideoCapture object to read from the video file.
    video_reader = cv2.VideoCapture(video_file_path)

    # Get the width and height of the video.
    original_video_width = int(video_reader.get(cv2.CAP_PROP_FRAME_WIDTH))
    original_video_height = int(video_reader.get(cv2.CAP_PROP_FRAME_HEIGHT))

    # Initialize the VideoWriter Object to store the output video in the disk.
    video_writer = cv2.VideoWriter(output_file_path, cv2.VideoWriter_fourcc('M', 'P', '4', 'V'),
                                   video_reader.get(cv2.CAP_PROP_FPS), (original_video_width,

    # Declare a queue to store video frames.
    frames_queue = deque(maxlen = SEQUENCE_LENGTH)

    # Initialize a variable to store the predicted action being performed in the video.
    predicted_class_name = ''

    # Iterate until the video is accessed successfully.
    while video_reader.isOpened():

        # Read the frame.
        ok, frame = video_reader.read()

        # Check if frame is not read properly then break the loop.
        if not ok:
            break

```

```

# Resize the Frame to fixed Dimensions.
resized_frame = cv2.resize(frame, (IMAGE_HEIGHT, IMAGE_WIDTH))

# Normalize the resized frame by dividing it with 255 so that each pixel value then 1
normalized_frame = resized_frame / 255

# Appending the pre-processed frame into the frames list.
frames_queue.append(normalized_frame)

# Check if the number of frames in the queue are equal to the fixed sequence length.
if len(frames_queue) == SEQUENCE_LENGTH:

    # Pass the normalized frames to the model and get the predicted probabilities.
    predicted_labels_probabilities = LRCN_model.predict(np.expand_dims(frames_queue,
    predicted_label = np.argmax(predicted_labels_probabilities)

    # Get the class name using the retrieved index.
    predicted_class_name = CLASSES_LIST[predicted_label]

# Write predicted class name on top of the frame.
cv2.putText(frame, predicted_class_name, (10, 30), cv2.FONT_HERSHEY_SIMPLEX, 1, (0, 2

# Write The frame into the disk using the VideoWriter Object.
video_writer.write(frame)

# Release the VideoCapture and VideoWriter objects.
video_reader.release()
video_writer.release()

```

```

1
2
3
4
5
6
7
8
# Construct the output video path.
output_video_file_path = f'{test_videos_directory}/{video_title}-Output-SeqLen{SEQUENCE_LENGTH}

# Perform Action Recognition on the Test Video.
predict_on_video(input_video_file_path, output_video_file_path, SEQUENCE_LENGTH)

# Display the output video.
VideoFileClip(output_video_file_path, audio=False, target_resolution=(300,None)).ipython_disp

```

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50


```

51
52
53
54
55
56
57
58
59
60
61
62
63
def predict_single_action(video_file_path, SEQUENCE_LENGTH):
    '''
    This function will perform single action recognition prediction on a video using the LRCN
    Args:
    video_file_path: The path of the video stored in the disk on which the action recognition
    SEQUENCE_LENGTH: The fixed number of frames of a video that can be passed to the model a
    '''

    # Initialize the VideoCapture object to read from the video file.
    video_reader = cv2.VideoCapture(video_file_path)

    # Get the width and height of the video.
    original_video_width = int(video_reader.get(cv2.CAP_PROP_FRAME_WIDTH))
    original_video_height = int(video_reader.get(cv2.CAP_PROP_FRAME_HEIGHT))

    # Declare a list to store video frames we will extract.
    frames_list = []

    # Initialize a variable to store the predicted action being performed in the video.
    predicted_class_name = ''

    # Get the number of frames in the video.
    video_frames_count = int(video_reader.get(cv2.CAP_PROP_FRAME_COUNT))

    # Calculate the interval after which frames will be added to the list.
    skip_frames_window = max(int(video_frames_count/SEQUENCE_LENGTH),1)

    # Iterating the number of times equal to the fixed length of sequence.
    for frame_counter in range(SEQUENCE_LENGTH):

        # Set the current frame position of the video.
        video_reader.set(cv2.CAP_PROP_POS_FRAMES, frame_counter * skip_frames_window)

        # Read a frame.
        success, frame = video_reader.read()

        # Check if frame is not read properly then break the loop.
        if not success:

```

```
break
```

```
# Resize the Frame to fixed Dimensions.
```

```
resized_frame = cv2.resize(frame, (IMAGE_HEIGHT, IMAGE_WIDTH))
```

```
# Normalize the resized frame by dividing it with 255 so that each pixel value then 1
```

```
normalized_frame = resized_frame / 255
```

```
# Appending the pre-processed frame into the frames list
```

```
frames_list.append(normalized_frame)
```

```
# Passing the pre-processed frames to the model and get the predicted probabilities.
```

```
predicted_labels_probabilities = convlstm_model.predict(np.expand_dims(frames_list, axis
```

```
# Get the index of class with highest probability.
```

```
predicted_label = np.argmax(predicted_labels_probabilities)
```

```
# Get the class name using the retrieved index.
```

```
predicted_class_name = CLASSES_LIST[predicted_label]
```

```
# Display the predicted action along with the prediction confidence.
```

```
print(f'Action Predicted: {predicted_class_name}\nConfidence: {predicted_labels_probabili
```

```
# Release the VideoCapture object.
```

```
video_reader.release()
```

```
1
```

```
2
```

```
3
```

```
4
```

```
5
```

```
6
```

```
7
```

```
8
```

```
9
```

```
10
```

```
11
```

```
# Download the youtube video.
```

```
video_title = download_youtube_videos('https://www.youtube.com/watch?v=XqqpZS0c1K0', test_vid
```

```
# Construct the input youtube video path
```

```
input_video_file_path = f'{test_videos_directory}/{video_title}.mp4'
```

```
# Perform Single Prediction on the Test Video.
```

```
predict_single_action(input_video_file_path, SEQUENCE_LENGTH)
```

```
# Display the input video.
```

```
VideoFileClip(input_video_file_path, audio=False, target_resolution=(300,None)).ipython_displ
```

